

GPU Computing

Lecture 2

Young-Ho Gong



광운대학교
KwangWoon University




Contents

- Graphics and Parallel Computing History
- 3D Graphics Pipeline
- Hardware Development
- Early GPGPU
- OpenGL
- GPU Shading Language
- CUDA
- Other Parallel Architectures



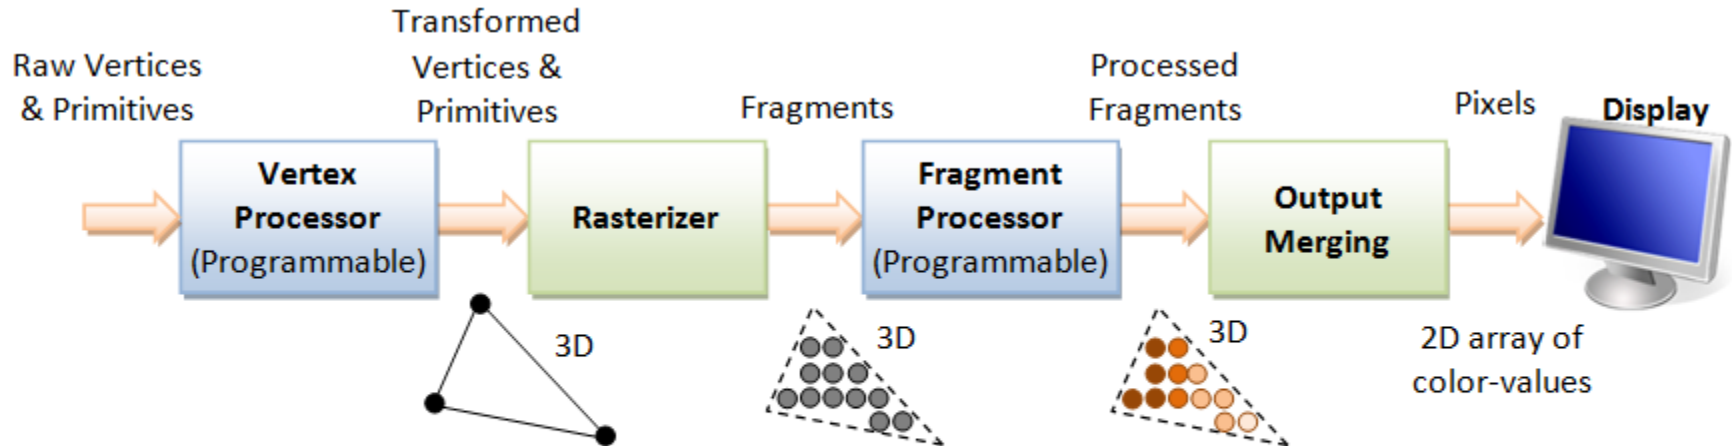
Computer Graphics

- Make great images
 - Intricate shapes
 - Complex optical effects
 - Seamless motion
- Make them **fast**
 - Invent clever techniques
 - Use every trick imaginable
 - ***Build monster hardware***
- We use the **monster hardware** as **computing devices**



Graphics Pipeline


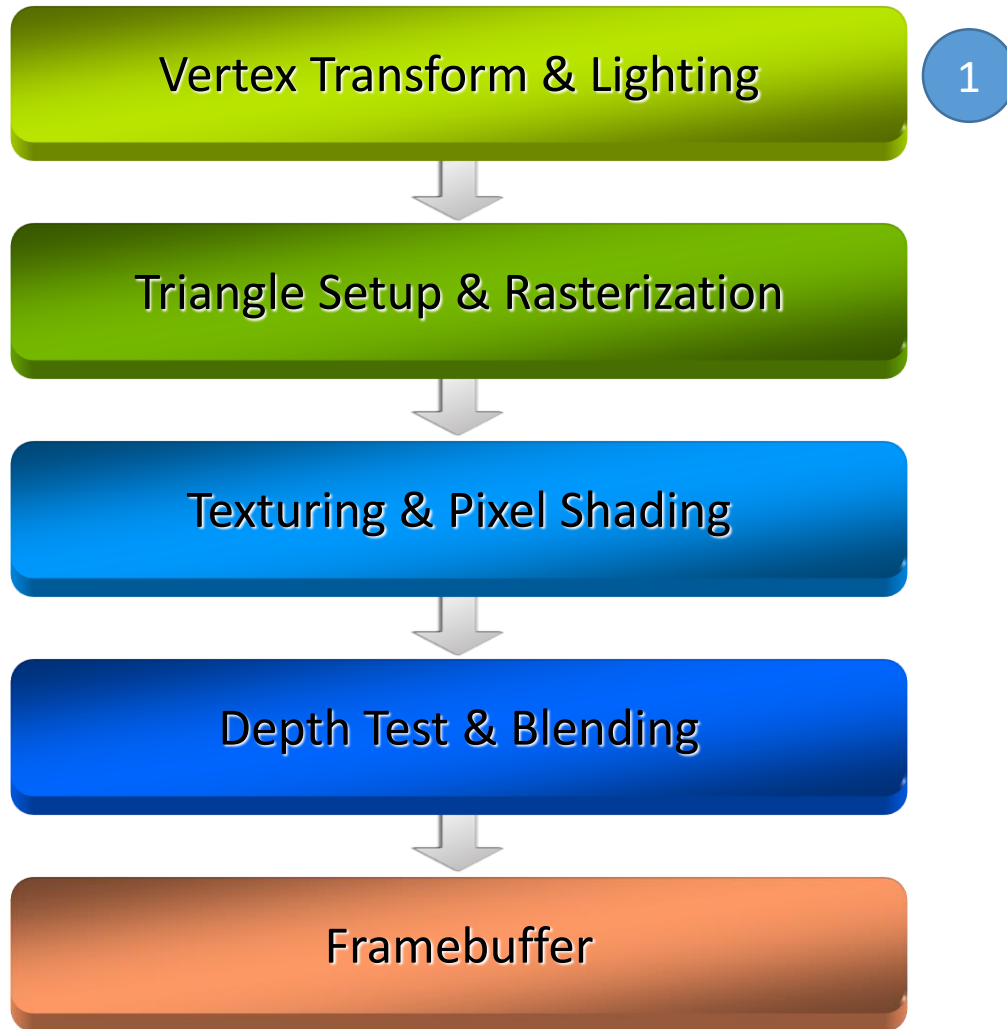
- Sequence of operations
 - take the vertices and textures of your meshes all the way to the pixels in the render targets




3D Graphics Rendering Pipeline: Output of one stage is fed as input of the next stage. A vertex has attributes such as (x, y, z) position, color (RGB or RGBA), vertex-normal (n_x, n_y, n_z) , and texture. A primitive is made up of one or more vertices. The rasterizer raster-scans each primitive to produce a set of grid-aligned fragments, by interpolating the vertices.




3D Graphics Pipeline




3D Graphics Pipeline




➔
Vertex shader




↙
Rasterization




➔
Fragment shader



3D Graphics Pipeline

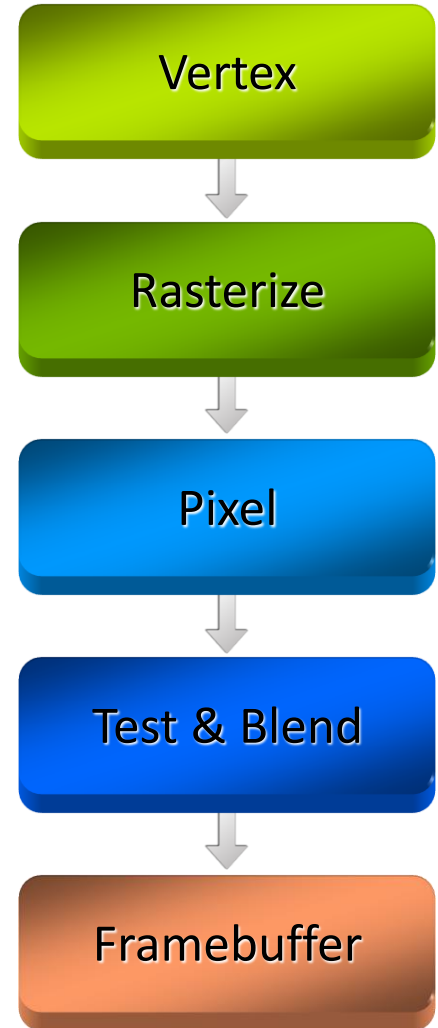


3D Graphics Pipeline



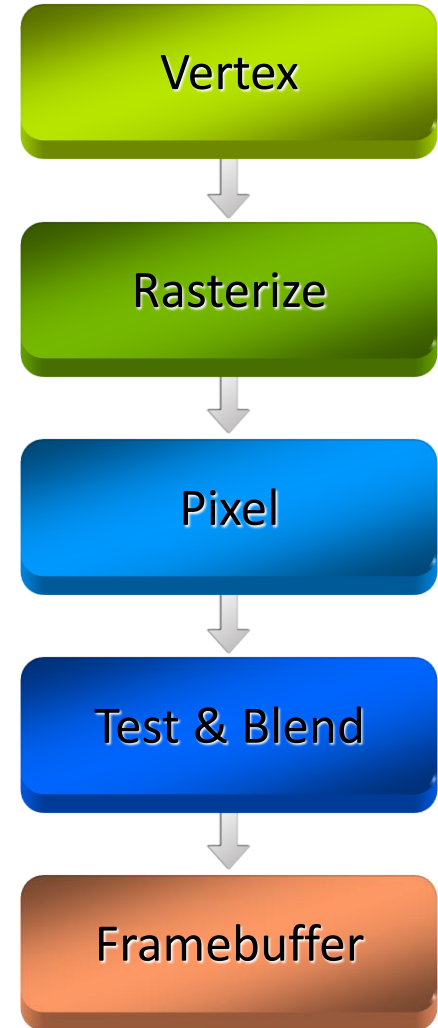
The Graphics Pipeline

- Key abstraction of real-time graphics
- One chip/board per stage
- Fixed data flow through pipeline



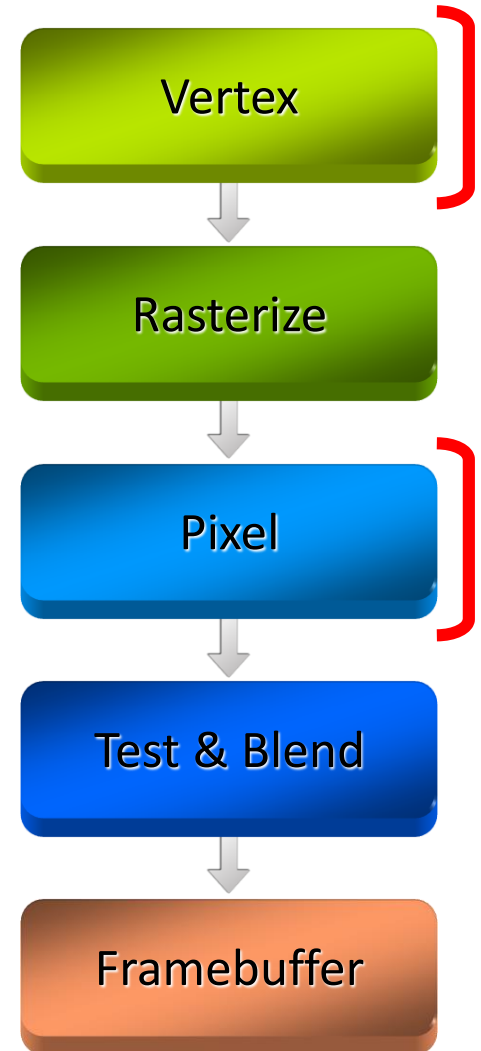
The Graphics Pipeline – 2nd gen

- Everything fixed function, with a certain number of modes
- Number of modes for each stage grew over time
- Hard to optimize HW
- Developers always wanted more flexibility



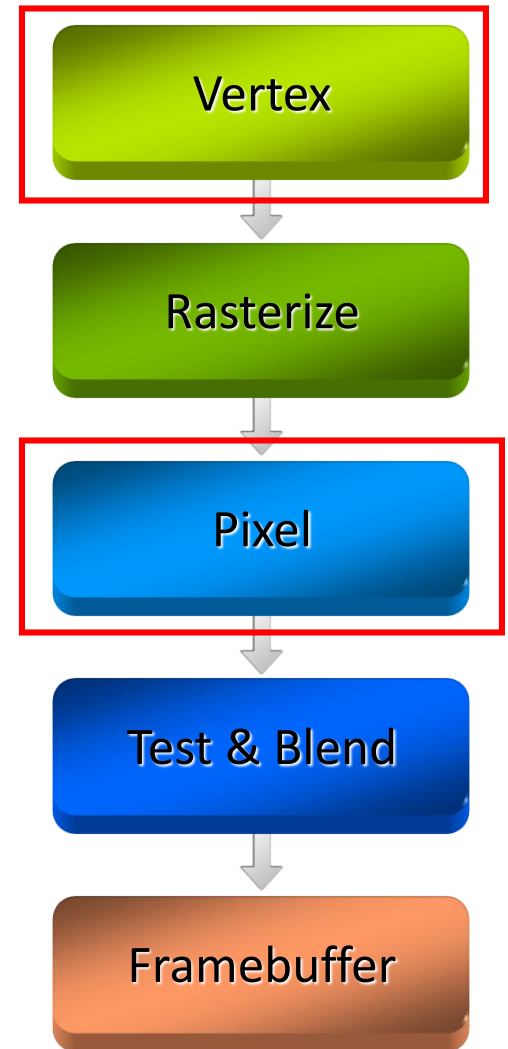
The Graphics Pipeline – 3rd gen

- Remains a key abstraction
- Vertex & pixel processing became programmable, new stages added
- GPU architecture increasingly centers around **shader** execution



The Graphics Pipeline – 4th gen

- Exposing a (at first limited) **instruction set** for some stages
- Expanded to **full ISA**
 - CUDA
- Limited instructions & instruction types and no control flow at first
- Optimization?




Why GPUs scale so nicely

- Workload and Programming Model provide **lots** of parallelism
- Applications provide **large groups of vertices at once**
 - Vertices can be processed **in parallel**
 - Apply same transform to all vertices
- Triangles contain **many pixels**
 - Pixels from a triangle can be processed **in parallel**
 - Apply same shader to all pixels
- **Very efficient hardware** to hide serialization bottlenecks



Why GPUs scale so nicely



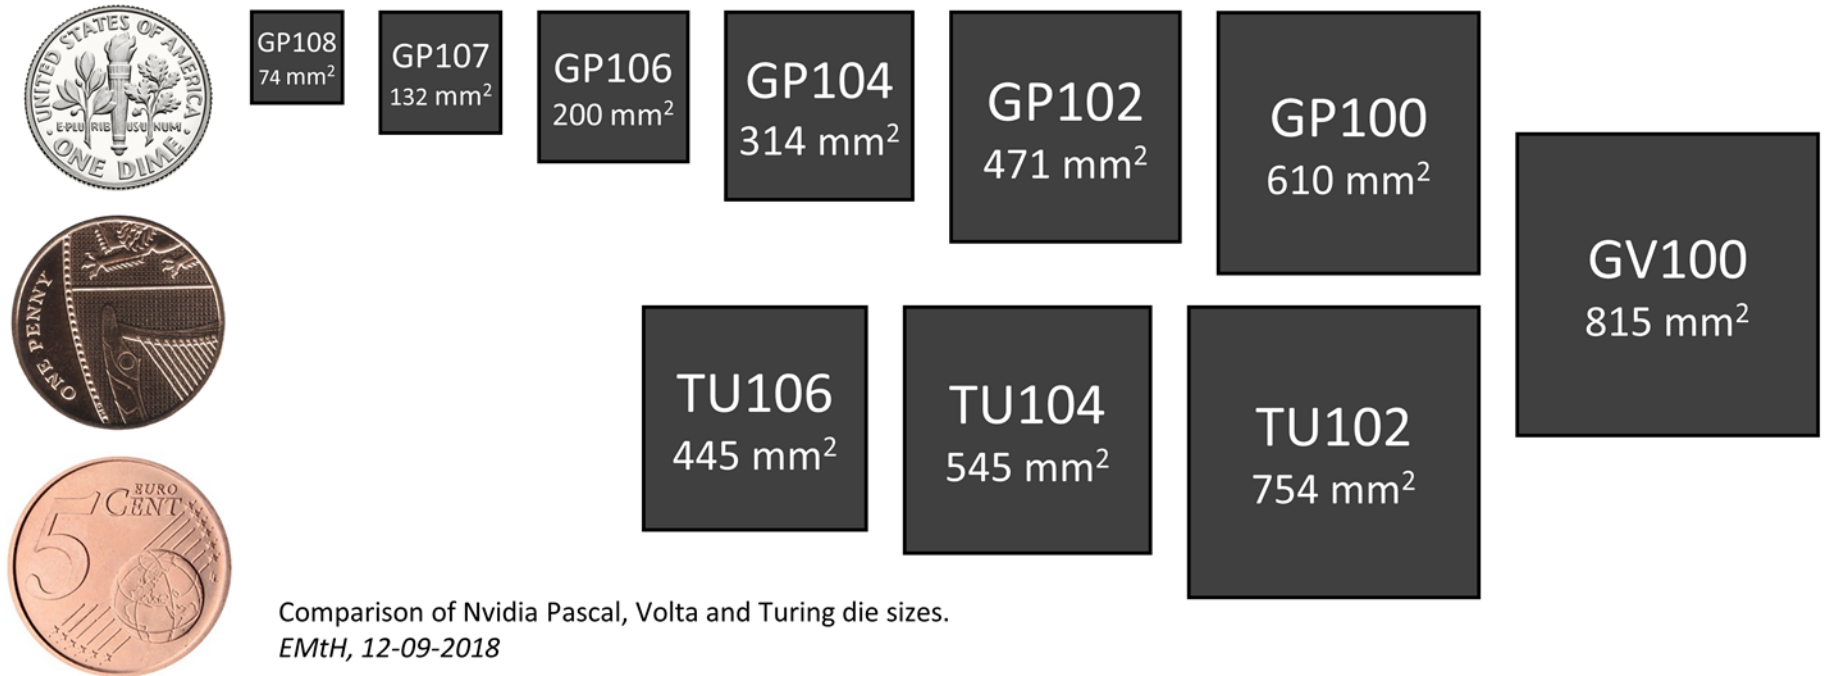
With Moore's Law...

- Old vs. New GPU



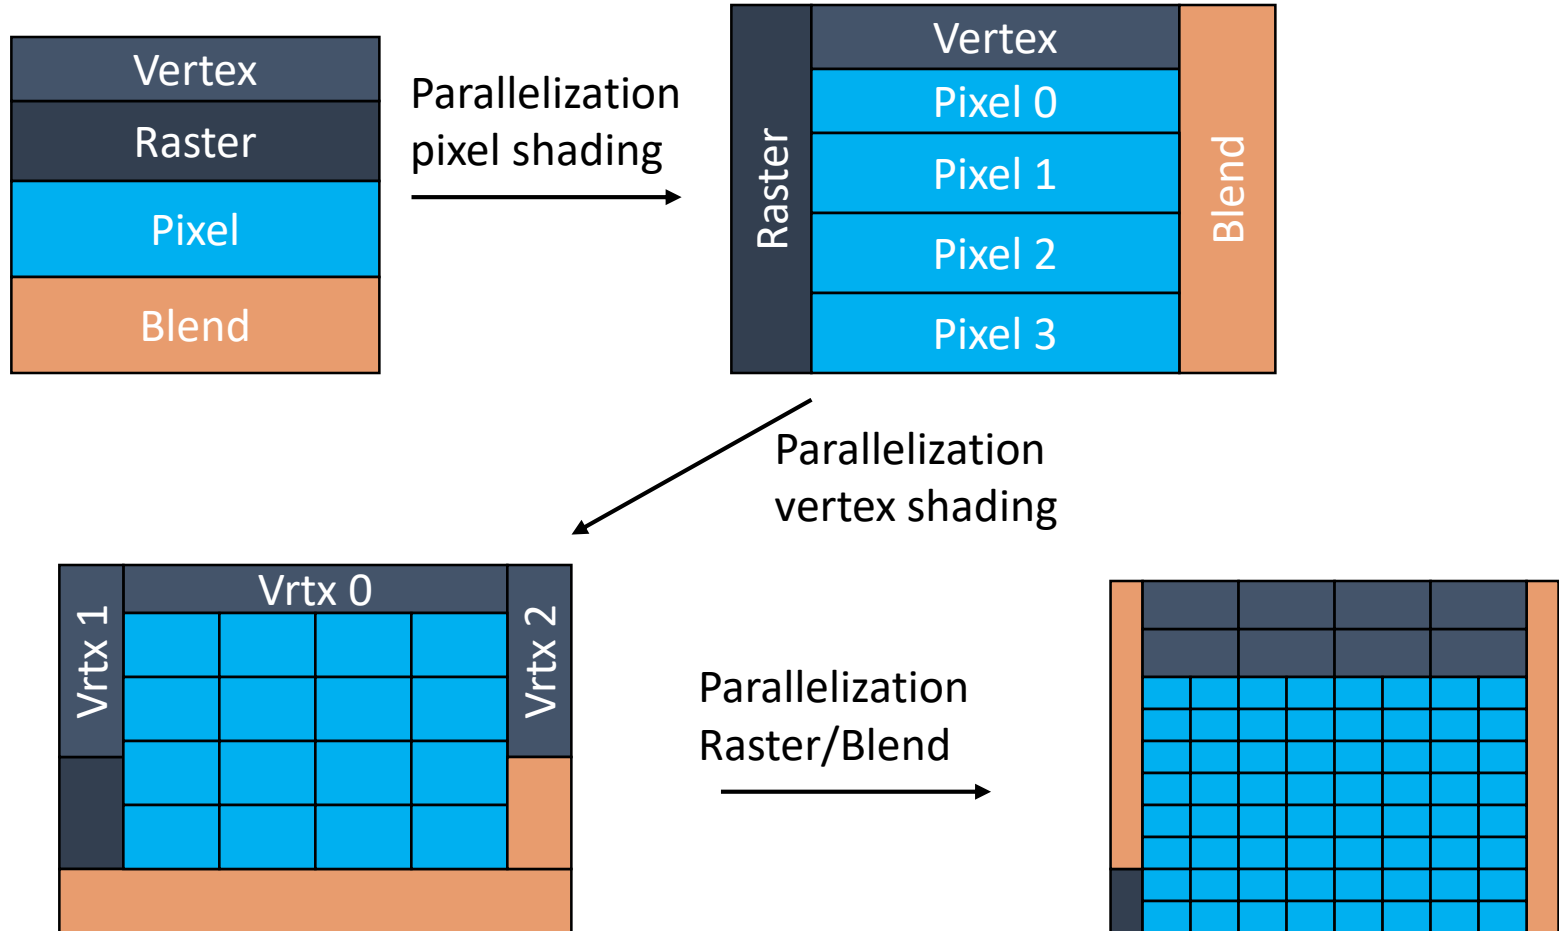
With Moore's Law...

- Low-end → High-end GPU



With Moore's Law...

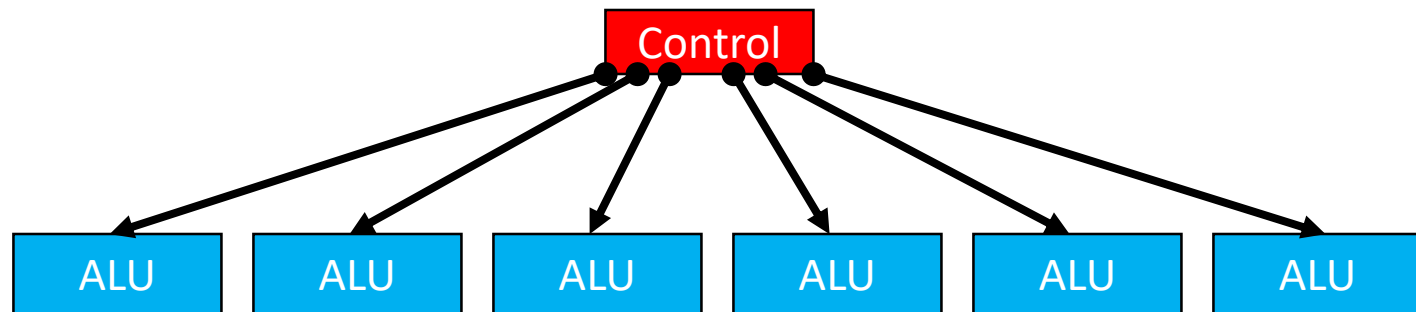

- We get more and more transistors !



More Efficiency

- Note that we do the **same thing** for lots of pixels/vertices

- In a CPU manner (Single control unit for each ALU)

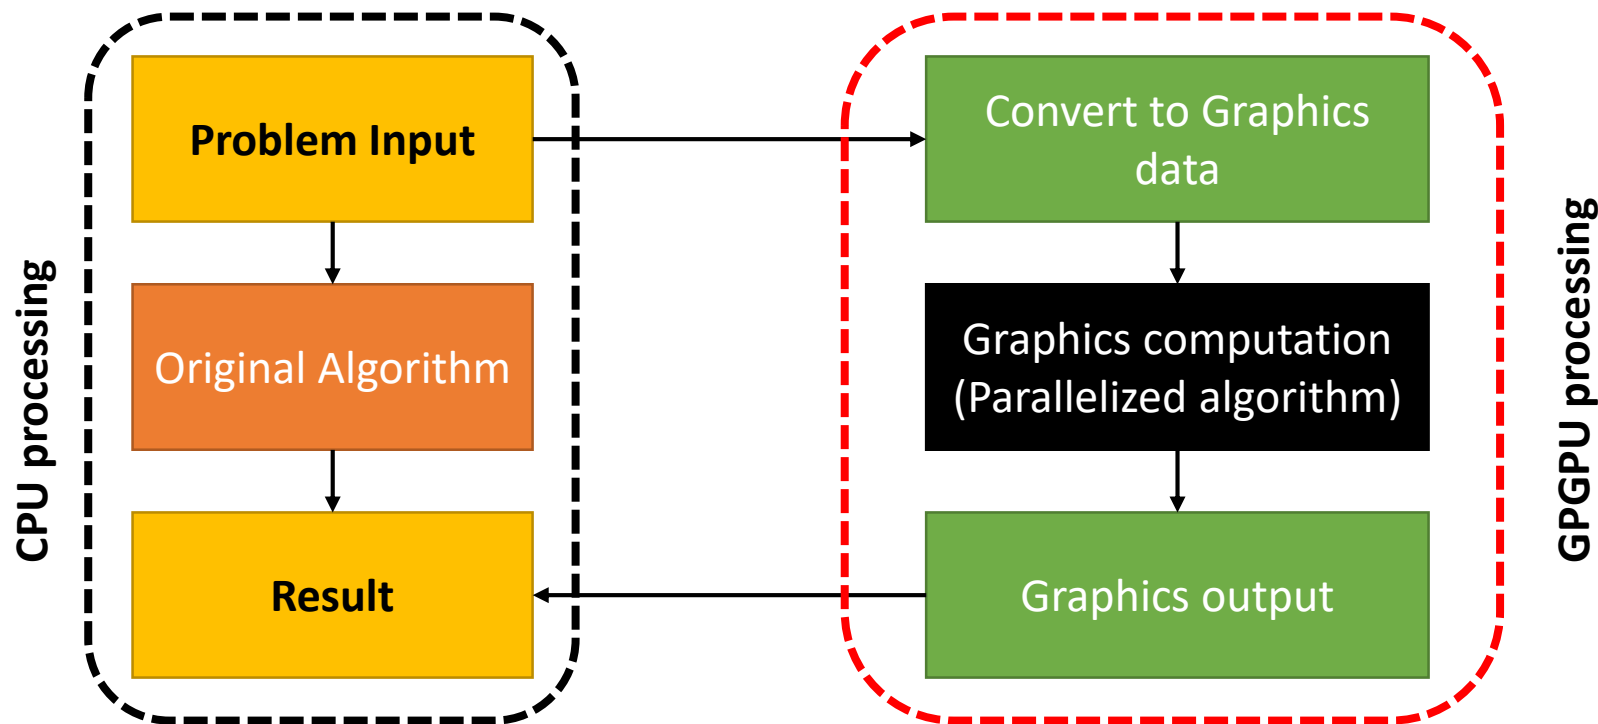


- A **warp** = 32 threads launched together
 - Usually, execute together as well



Early GPGPU

- GPGPU: General Purpose Graphics Processing Unit



- Very tedious, **limited usability**
- Still had some very nice results (Performance)
- This was the lead up to **CUDA**



CUDA

- Compute Unified Device Architecture
- General purpose programming model
 - User kicks off batches of threads on the GPU
 - GPU = dedicated super-threaded, massively data parallel co-processor
- Targeted software stack
 - Compute oriented drivers, language, and tools
- Driver for loading computation programs into GPU
 - Standalone Driver - Optimized for computation
 - Interface designed for compute - graphics free API
 - Data sharing with OpenGL buffer objects
 - Guaranteed maximum download & readback speeds
 - Explicit GPU memory management



CUDA

■ Compute Unified Device Architecture

GPU Computing Applications						
Libraries and Middleware						
cuDNN TensorRT	cuFFT, cuBLAS, cuRAND, cuSPARSE	CULA MAGMA	Thrust NPP	VSIPL, SVM, OpenCurrent	PhysX, OptiX, iRay	MATLAB Mathematica
Programming Languages						
C	C++	Fortran	Java, Python, Wrappers	DirectCompute	Directives (e.g., OpenACC)	
CUDA-enabled NVIDIA GPUs						
Turing Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier	GeForce 2000 Series	Quadro RTX Series	Tesla T Series		
Volta Architecture (Compute capabilities 7.x)	DRIVE/JETSON AGX Xavier			Tesla V Series		
Pascal Architecture (Compute capabilities 6.x)	Tegra X2	GeForce 1000 Series	Quadro P Series	Tesla P Series		
Maxwell Architecture (Compute capabilities 5.x)	Tegra X1	GeForce 900 Series	Quadro M Series	Tesla M Series		
Kepler Architecture (Compute capabilities 3.x)	Tegra K1	GeForce 700 Series GeForce 600 Series	Quadro K Series	Tesla K Series		
	EMBEDDED	CONSUMER DESKTOP, LAPTOP	PROFESSIONAL WORKSTATION	DATA CENTER		



Other Architectures

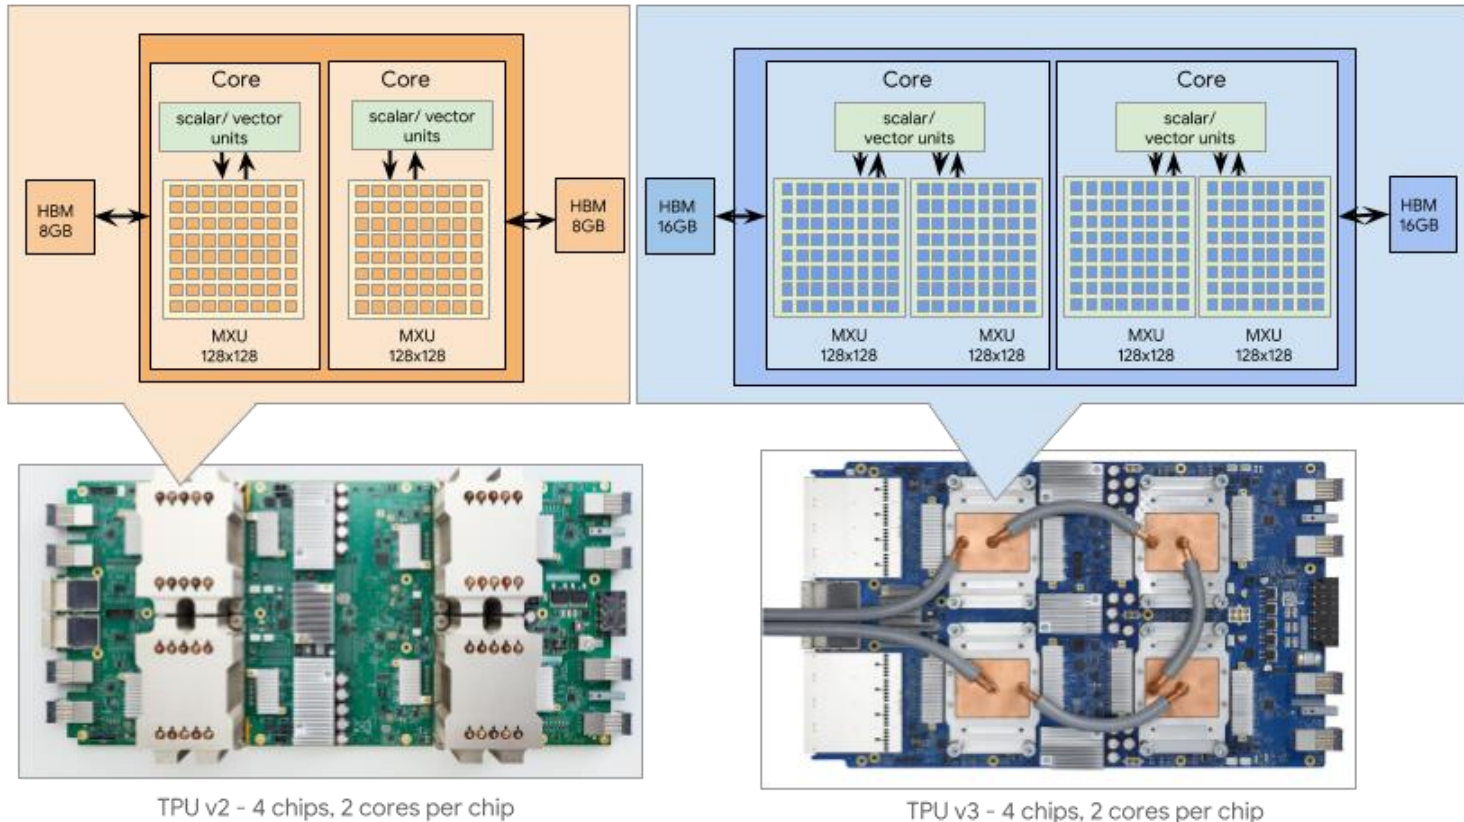
- IBM Watson Architecture
 - Traditional design
 - 90 * IBM Power 750 system (3.5GHz IBM Power7 8-core CPU)
 - Each core can execute 4 threads $\rightarrow 90 * 8 * 4 = 2,880$ threads totally
 - 2,880 threads + 16 TB RAM
- It works as a question answering computing system
 - w/ Deep learning technology



Other Architectures

■ AlphaGo Architecture

- Deep Learning → so much computations
- Cloud computing required → scalable architecture
- Oct 2015: 1202 CPU + 176 GPU (CUDA) – **fake**
- Mar 2016: 1920 CPU + 280 GPU (CUDA) – **fake**
- May 2016: AlphaGo uses **TPU** !



Other Architectures

■ AI

Scalar

Vector

Matrix

Tensor


1

$$\begin{bmatrix} 1 \\ 2 \end{bmatrix}$$
$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$
$$\begin{bmatrix} \begin{bmatrix} 1 & 2 \end{bmatrix} & \begin{bmatrix} 3 & 2 \end{bmatrix} \\ \begin{bmatrix} 1 & 7 \end{bmatrix} & \begin{bmatrix} 5 & 4 \end{bmatrix} \end{bmatrix}$$


vector




matrix




3d-tensor



4d-tensor



5d-tensor



6d-tensor



Intermediate Summary

- Graphics and Parallel Computing History
- 3D Graphics Pipeline
- HW Development
- Early GPGPU
- CUDA
- Other Parallel Architectures



Next step

- CUDA Install in your computer
- If your desktop/laptop has an Nvidia GPU
 - Install Visual Studio 2019
 - Install latest CUDA SDK
 - You can start CUDA programming by generating new CUDA project
- Else,
 - You can start CUDA Programming using Colab by Google



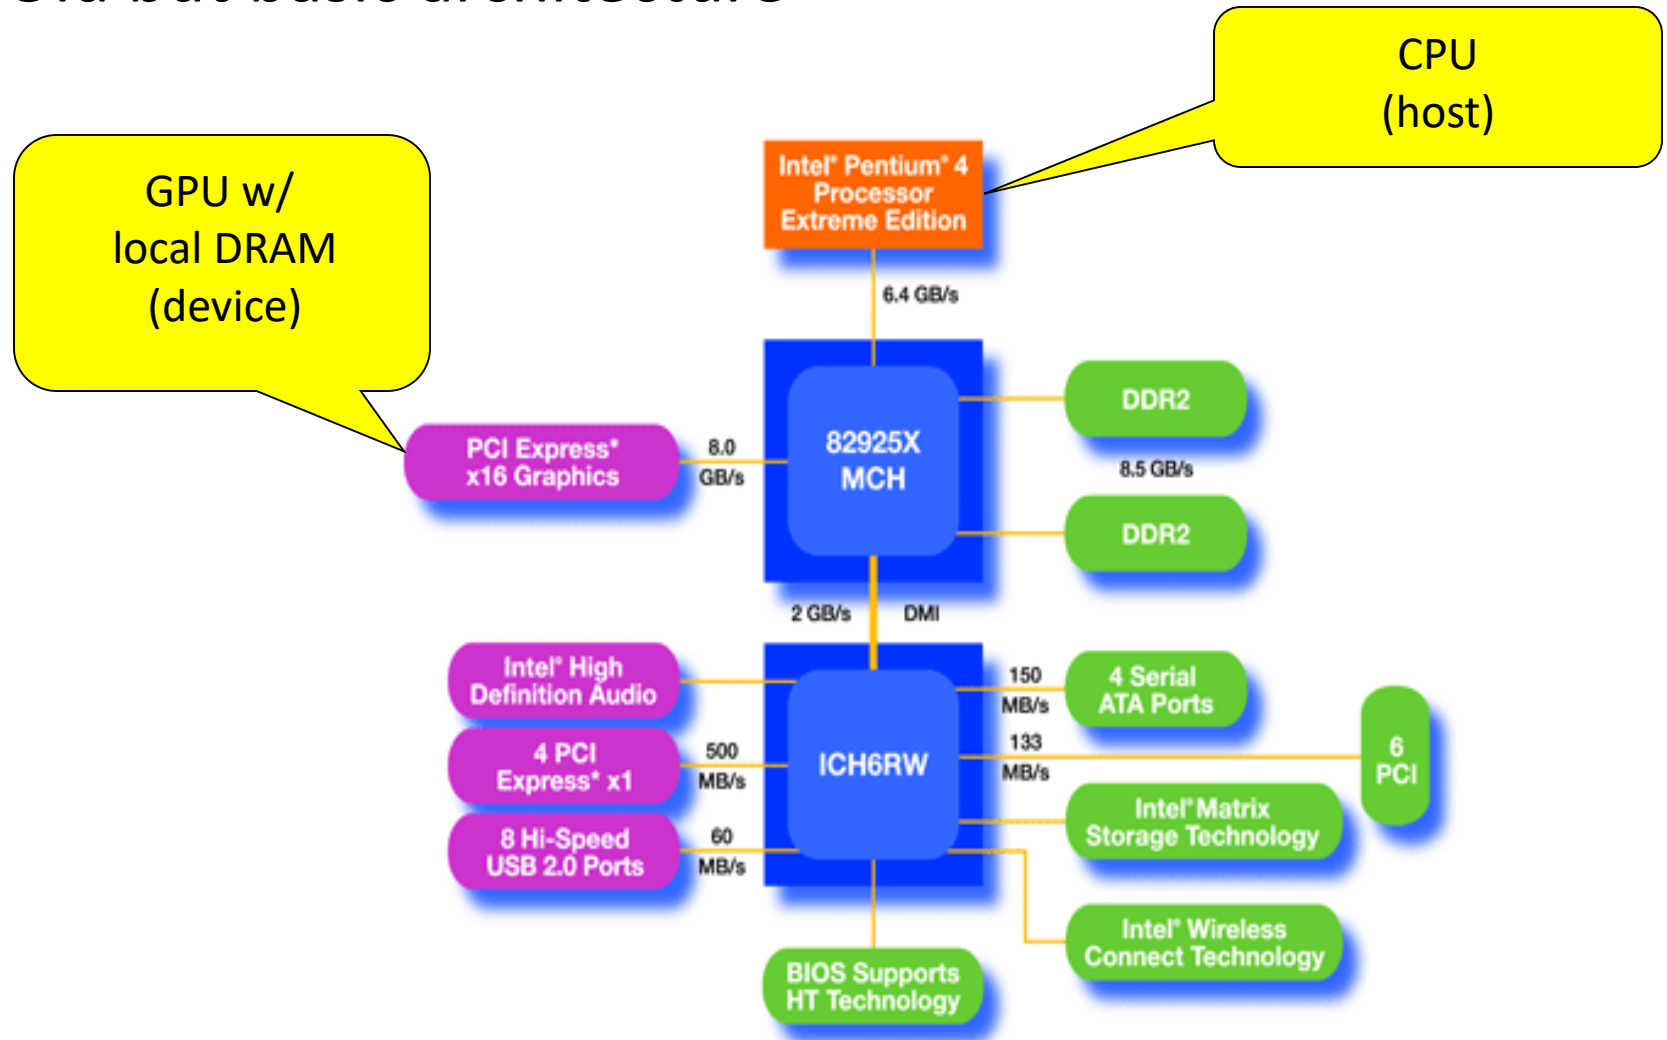
Next step

- Very simple CUDA Program
- Simple CUDA model
 - Host, device
- CUDA memory copy scenario
- CPU memory copy functions
- CUDA memory copy functions
- Host-device memory copy functions




An Example of Physical Reality behind CUDA

- Old but basic architecture



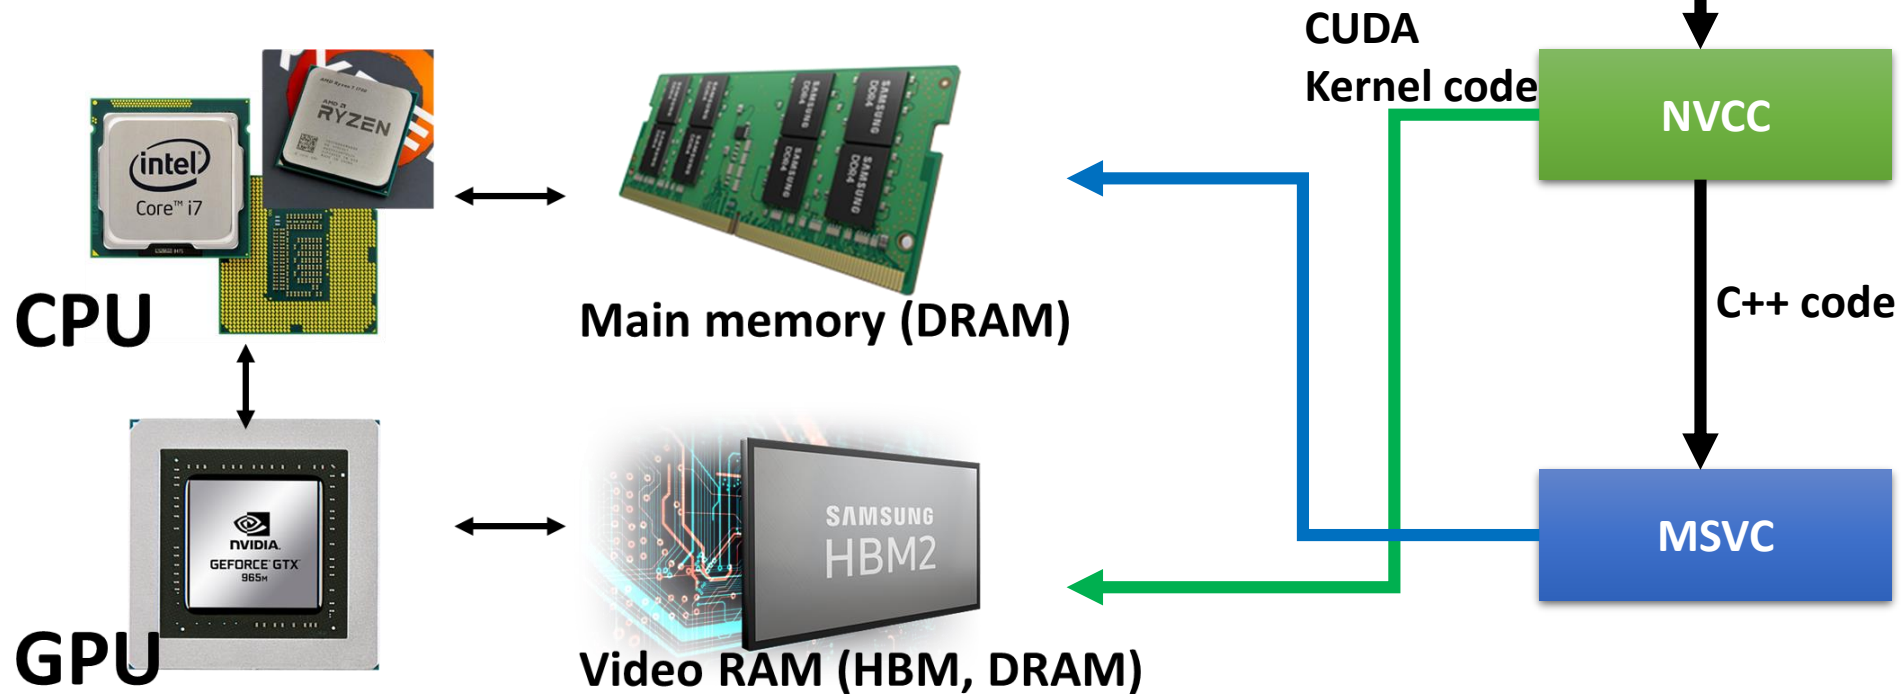
Very simple CUDA model

- **Host:** CPU + main memory
- **Device:** GPU + VRAM



Very simple CUDA model

- nvcc: NVIDIA CUDA compiler
- msvc: Microsoft visual C++ compiler



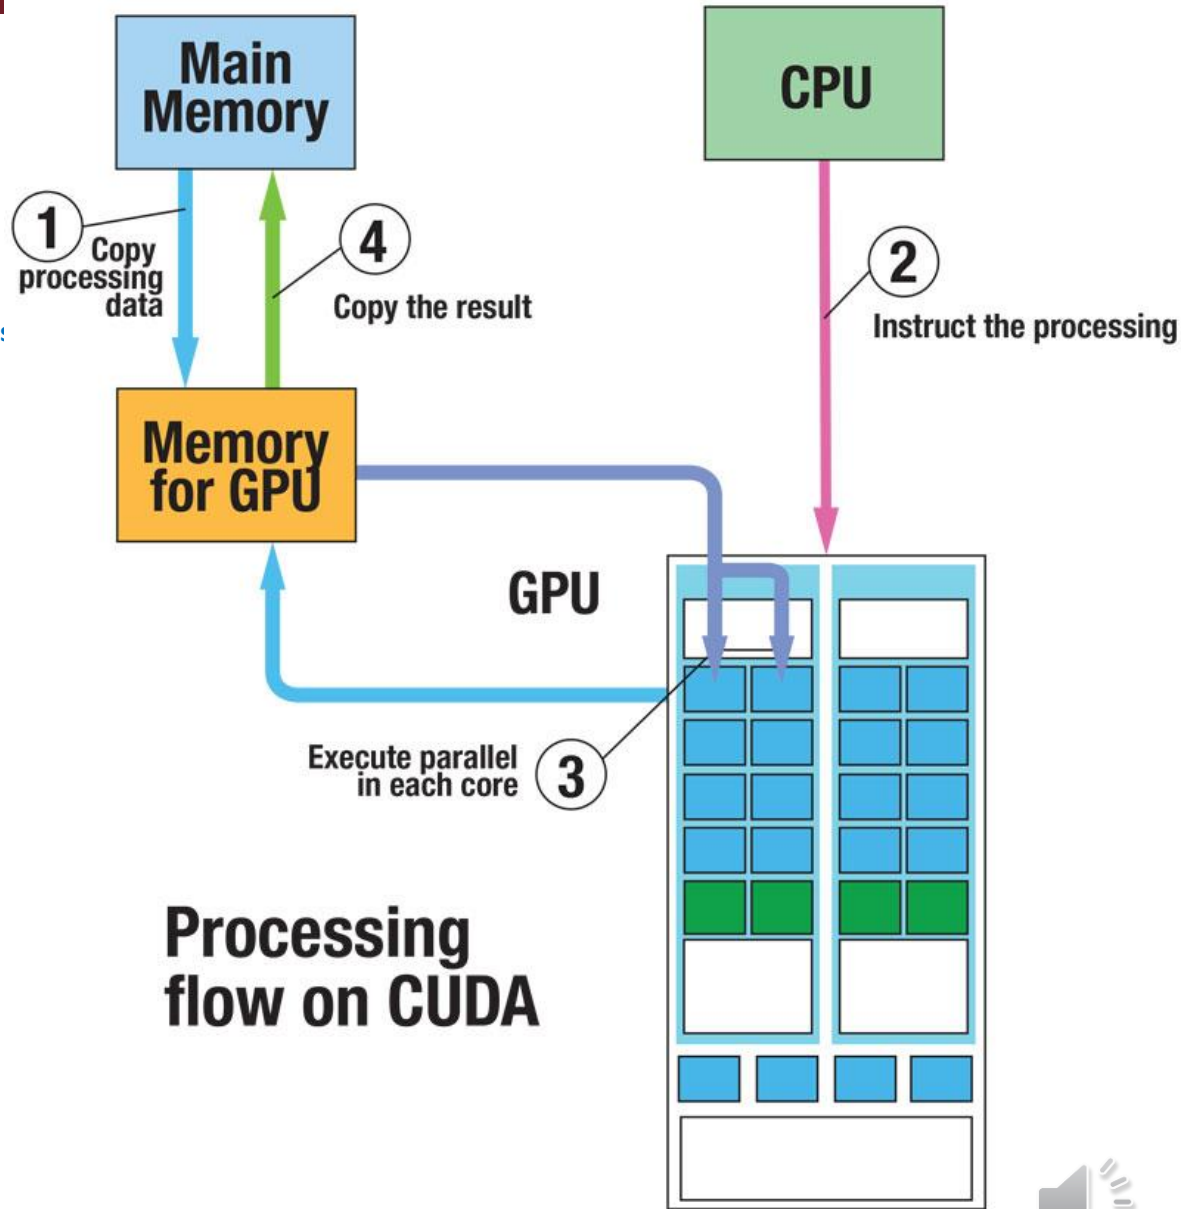
Simple CUDA Program Scenario

- Step 1. Serial code (Host)
 - Serial execution: **read data**
 - Prepare parallel execution
 - Copy data from main memory to VRAM
- Step 2. Kernel code (Device)
 - **Parallel Processing**
 - Read/Write data from VRAM to VRAM
 - **CUDA do NOT have any direct I/O features**
- Step 3. Serial code (Host)
 - Copy data from VRAM to main memory
 - Serial execution: **print data**



Simple CUDA Program Scenario

- Step 1. Serial code (Host)
 - Serial execution: **read data**
 - Prepare parallel execution
 - Copy data from main memory to VRAM
- Step 2. Kernel code (Device)
 - **Parallel Processing**
 - Read/Write data from VRAM to VRAM
 - **CUDA do NOT have any direct I/O feature!**
- Step 3. Serial code (Host)
 - Copy data from VRAM to main memory
 - Serial execution: **print data**



Memory Spaces

- CPU and GPU have separate memory spaces
 - Data is moved across data bus (e.g., PCIe)
 - Use functions to allocate/set/copy memory on GPU
 - Very similar to corresponding C functions
- Pointers are just addresses
 - Can't tell from the pointer value whether the address is on CPU or GPU
 - Must exercise care when dereferencing:
 - Dereferencing CPU pointer on GPU will likely crash
 - Same for vice versa



CPU Memory Allocation / Release

- Host (CPU) manages host (CPU) memory:
 - void* malloc (size_t nbytes)
 - void* memset (void * pointer, int value, size_t count)
 - void free (void* pointer)

```
int n = 1024;
```

```
int nbytes = 1024*sizeof(int);
```

```
int * ptr = 0;
```

```
ptr = malloc(nbytes);
```

```
memset( ptr, 0, nbytes);
```

```
free(ptr);
```



GPU Memory Allocation / Release

- Host (CPU) manages device (GPU) memory:
 - cudaMalloc (**void ** pointer**, size_t nbytes)
 - cudaMemcpy (void * pointer, int value, size_t count)
 - cudaFree (void* pointer)

```
int n = 1024;
```

```
int nbytes = 1024*sizeof(int);
```

```
int * dev_a = 0;
```

```
cudaMalloc( (void**)&dev_a, nbytes );
```

```
cudaMemset( dev_a, 0, nbytes);
```

```
cudaFree(dev_a);
```



CUDA function rules

- Every library function starts with “**cuda**”
- Most of them returns error code (or **cudaSuccess**).
 - `cudaError_t cudaMalloc(void** devPtr, size_t size);`
 - `cudaError_t cudaFree(void* devPtr);`
 - `cudaError_t cudaMemcpy(void* dst, const void* src, size_t size);`
- Example:
 - ```
if (cudaMalloc(&devPtr, SIZE) != cudaSuccess){
 exit(1);
}
```



# Data copies

- `cudaError_t cudaMemcpy (void* dst ,  
                                void* src,  
                                size_t nbytes,  
                                enum cudaMemcpyKind direction);`
  - Returns after the copy is complete
  - Blocks CPU thread until all bytes have been copied
  - Doesn't start copying until previous CUDA calls complete
- `enum cudaMemcpyKind`
  - `cudaMemcpyHostToDevice`
  - `cudaMemcpyDeviceToHost`
  - `cudaMemcpyDeviceToDevice`
  - `cudaMemcpyHostToHost` (added !)



# Assignment #1

---

- Cuda Malloc
  - Host to device
  - Device to host



# Summary

---

- Very simple CUDA program
- Simple CUDA model
  - Host, Device
- CUDA memory copy scenario
- CPU memory copy functions
- CUDA memory copy functions
- Host-device memory copy example



# Thank you

Any questions?

E-mail: [yhgong@kw.ac.kr](mailto:yhgong@kw.ac.kr)  
Lab: <https://sites.google.com/view/yhgong/>

