

GPU Computing

Lecture 3

Young-Ho Gong



광운대학교
KwangWoon University



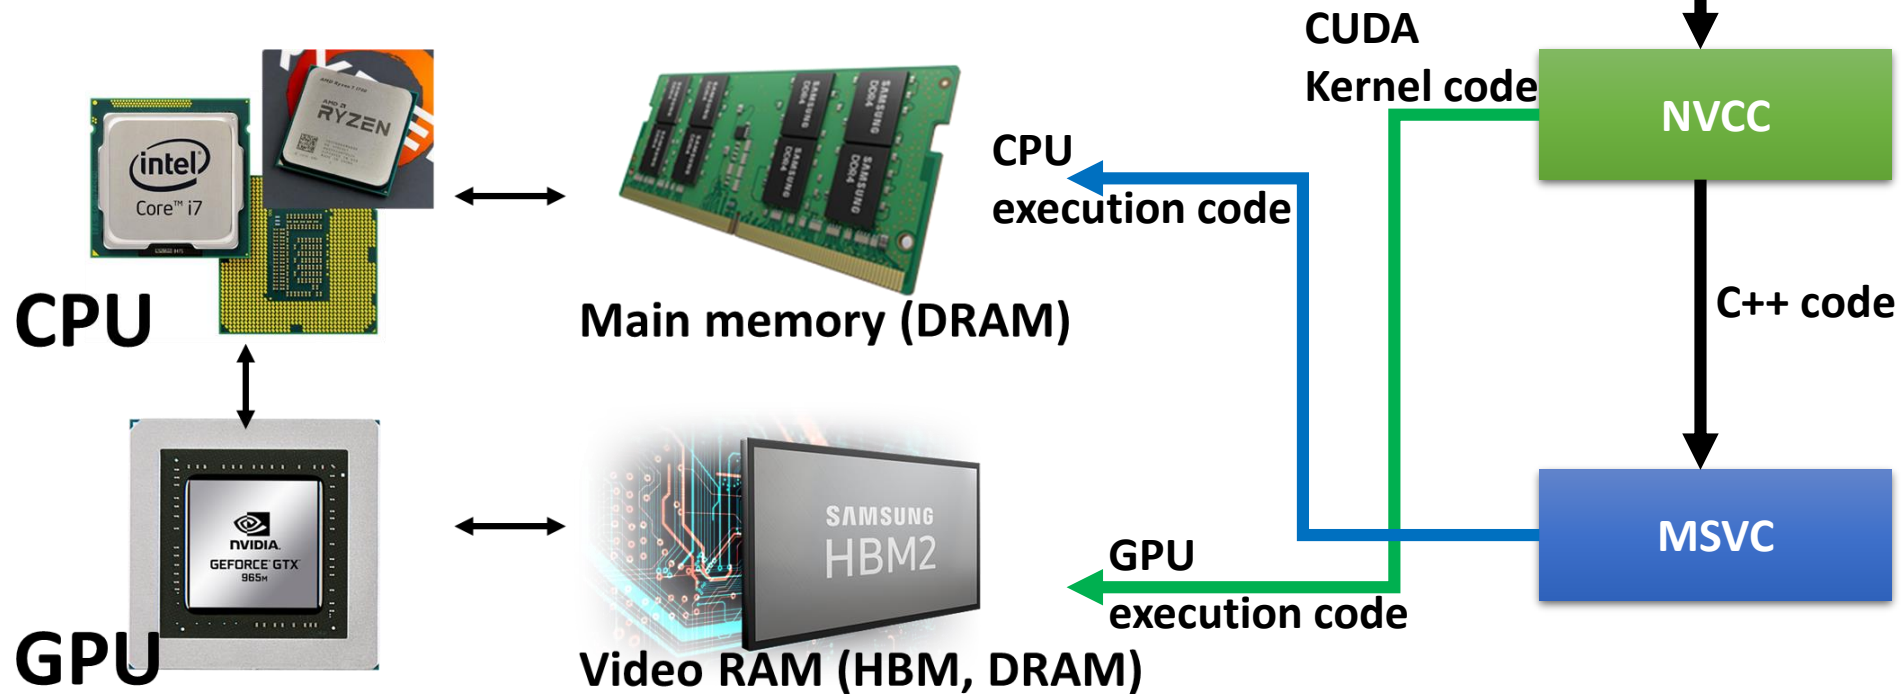
Contents

- Kernel concept
 - the central or most important part of something.
 - "this is the kernel of the argument"
- CUDA programming model
- CUDA function declarations
- Vector addition example
- CPU implementation
 - For loop
- Kernel concept
 - Function, as loop-body
- CPU kernel implementation
- CUDA implementation
 - Multiple thread launch
- Kernel launch
 - Example code



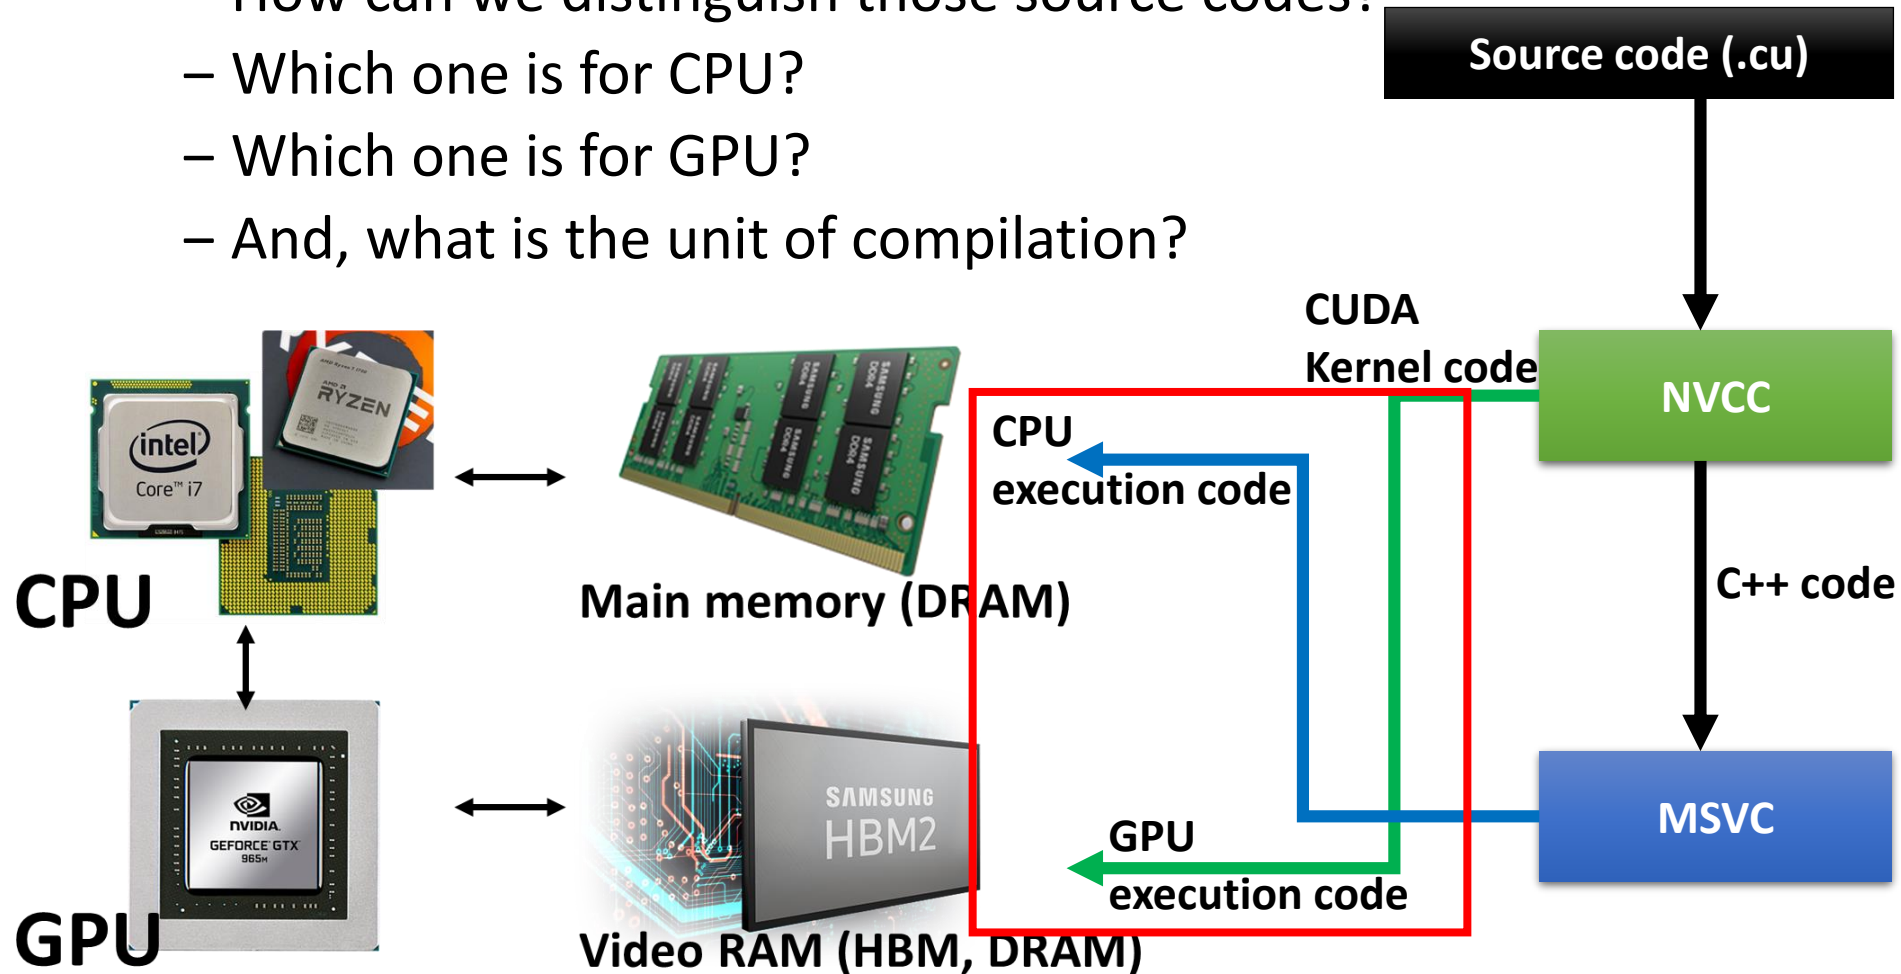
CUDA programming model

- We have two sets of execution code
 - For CPU (host part), msvc is used
 - For GPU (device part), nvcc is used



CUDA programming model

- At the source code level,
 - How can we distinguish those source codes?
 - Which one is for CPU?
 - Which one is for GPU?
 - And, what is the unit of compilation?



CUDA programming model

- Compilation unit?

- Function
- Each function will be assigned to CPU and/or GPU

- How to distinguish them?

- Use PREFIX for each function

`__host__` : can be called by CPU (default, can be omitted)

`__device__` : called from other GPU functions,
cannot be called by the CPU

`__global__` : launched by CPU, (bridge between host/device)
cannot be called from GPU, must return void

`__host__` and `__device__` qualifiers can be combined



More on CUDA Function Declarations

- `__global__` defines a kernel function
 - Each “`__`” consists of two underscore characters
 - A kernel function must return void
- `__device__` and `__host__` can be used together
 - Compiled twice (!)

	Executed on the:	Only callable from the:
<code>__device__ float DeviceFunc()</code>	Device	Device
<code>__global__ void KernelFunc()</code>	Device	Host
<code>__host__ float HostFunc()</code>	Host	Host



Example Source Code

```
__device__ inline void myAtomicADD(unsigned long long int* addr){
    unsigned long long int oldVal = *address;
    unsigned long long int newVal = oldVal + val;
    unsigned long long int readback;
    while ((readback = atomicCAS(address, oldVal, newVal)) != oldVal){
        oldVal = readback;
        newVal = oldVal + val;
    }
}

__global__ void kernel(unsigned long long int* pCount){
    myAtomicAdd(pCount, 1ULL);
}

__host__ int main(void){
    unsigned long long int aCount[1];
    ...
}
```



GPU Executable Functions

- CUDA language == C/C++ with some restrictions:
 - Can only access **GPU memory**
 - In new versions, can access host memory directly, with performance drawbacks.
 - **No recursion**
 - **No dynamic polymorphism**
 - No variable number of arguments
 - No static variables
 - No static variable declarations inside the function



Example: Hello.cu

- Make a new source code, hello.cu
- .cu file = .cpp file + CUDA kernel source code
 - Default function: `__host__`
- Hello.cu
 - `#include <stdio.h>`

```
int main (){  
    printf( "Hello, world!\n" );  
    return 0;  
}
```

- Result
 - Hello, world!



Programming TIP

- Sometimes CUDA does not output anything from `printf()`.
- Try to use `fflush()`
 - `Printf` → `stdio` buffer → output
 - `Fflush`: flush `stdio` buffer

- `Hello.cu`
 - `#include <stdio.h>`

```
int main (){  
    printf( "Hello, world!\n" );  
    fflush (stdout);  
    return 0;  
}
```




Scenario: Vector Addition


- We use arrays for vector representation

- $C[i] = A[i] + B[i]$

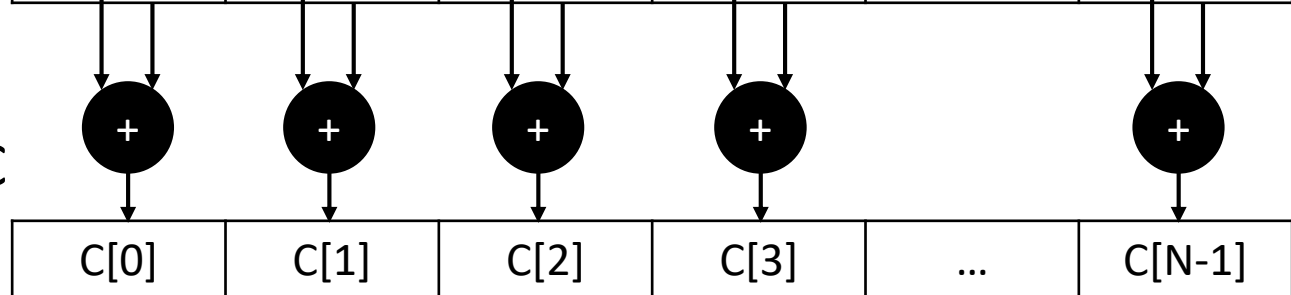
- Vector A



- Vector B



- Vector C



Vector Addition

- Vector: represented as 1D array
 - `const int a[SIZE];`
 - `const int b[SIZE];`
 - `int c[SIZE];`
- Vector addition: $c[\dots] = a[\dots] + b[\dots]$
- Serial execution: for-loop
 - E.g., `for(i=0;i<SIZE;i++)`
`c[i]=a[i]+b[i];`
- **A single CPU** does for $i=0,1,2,\dots,SIZE-1$
- CUDA execution: parallel kernel execution



Introducing a kernel function

- Converting to a kernel function,
void addKernel(int* dst, int aVal, int bVal){
 ***dst = aVal + bVal;**
}
...
- In main function,

for (int i=0;i<SIZE;i++)
 addKernel(&c[i],a[i],b[i]);
}
- **Kernel function:**
 - Substitute for the **loop body**
 - With proper data values



Introducing a kernel function

- Another view
 - CPU[0] executes `addKernel()` with its own data
 - CPU[1] executes `addKernel()` with its own data
 - CPU[2] executes `addKernel()` with its own data
 - ...
 - CPU[SIZE-1] executes `addKernel()` with its own data
- We have a single CPU → sequential execution
- If we have many CPUs? → Parallel execution possible!



CUDA kernel

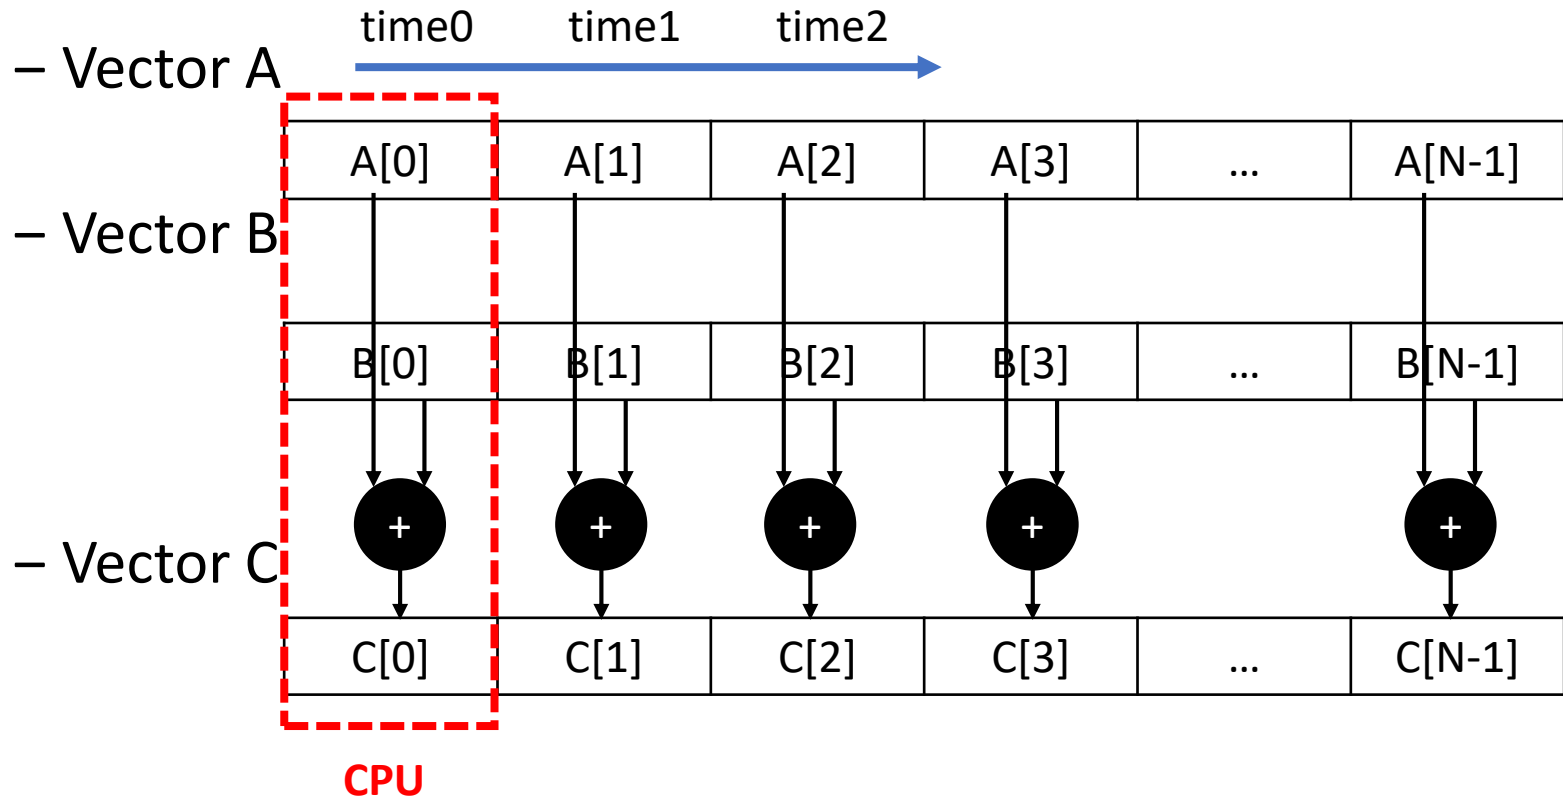
- CPU kernels
 - With a single CPU
 - For-loop
- Sequential execution
- For-loop
 - CPU[0]
 - CPU[1]
 - CPU[2]
 - ...
 - CPU[N-1]
- GPU kernels
 - A set of GPU cores
 - Multiple threads
- Parallel execution
- Kernel launch
 - GPU[0]
 - GPU[1]
 - GPU[2]
 - ...
 - GPU[N-1]



CPU-based vector addition

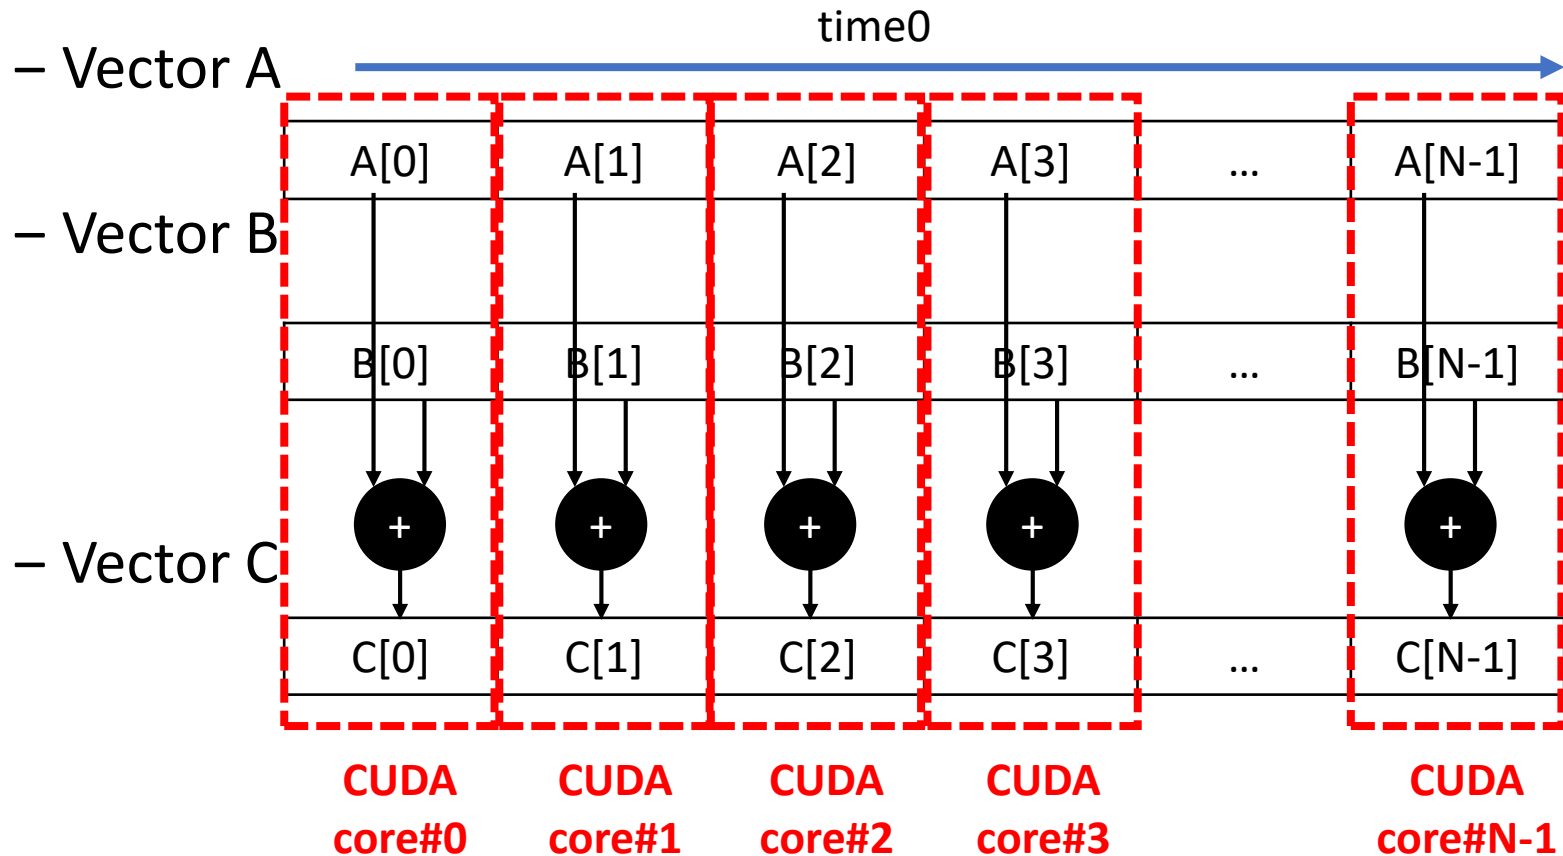
- A single CPU does an addition

- $C[i] = A[i] + B[i]$




CUDA-based vector addition

- We have many GPU cores
 - They do the addition at the same time!



Scenario: CUDA vector addition

- Step 1: Host-side
 - Make A, B with source data
 - Prepare C for the result
- Step 2: data copy host → device
 - cudaMemcpy from host to device
- Step 3: addition in CUDA
 - **Kernel launch for CUDA device**
 - Result will be stored in device memory (VRAM)
- Step 4: data copy device → host
 - cudaMemcpy from device to host
- Step 5: host-side
 - Print the output data



CUDA kernel launch syntax

- Prepare a CUDA kernel function
 - `__global__ void addKernel(int *c, const int* a, const int* b){
 int i = threadIdx.x;
 c[i] = a[i] + b[i];
}`
- Kernel launch syntax
 - `addKernel <<<1,SIZE>>>(dev_c, dev_a, dev_b);`
 - `Function_Name<<<# of Blocks,# of Threads>>>(arg1,arg2,...);`
- CUDA view
 - A thread executes `addKernel()` with `threadIdx.x = 0`
 - A thread executes `addKernel()` with `threadIdx.x = 1`
 - A thread executes `addKernel()` with `threadIdx.x = 2`
 - ...
 - A thread executes `addKernel()` with `threadIdx.x = SIZE-1`



Example code

- Kernel function



```
%%CU
```

```
#include <stdio.h>
```

```
__global__ void addKernel(int *c, const int *a, const int *b)  
{  
    int i = threadIdx.x;  
    c[i] = a[i] + b[i];  
}
```



Example code

```
int main()
{
    const int SIZE = 5;
    const int a[SIZE] = { 1, 2, 3, 4, 5 };
    const int b[SIZE] = { 10, 20, 30, 40, 50 };
    int c[SIZE] = { 0 };
```

```
int *dev_a = 0;
int *dev_b = 0;
int *dev_c = 0;

cudaMalloc((void**)&dev_c, SIZE * sizeof(int));
cudaMalloc((void**)&dev_a, SIZE * sizeof(int));
cudaMalloc((void**)&dev_b, SIZE * sizeof(int));
```



Example code

```
cudaMemcpy(dev_a, a, SIZE * sizeof(int), cudaMemcpyHostToDevice);  
cudaMemcpy(dev_b, b, SIZE * sizeof(int), cudaMemcpyHostToDevice);  
addKernel<<<1, SIZE>>>(dev_c, dev_a, dev_b);
```

```
cudaMemcpy(c, dev_c, SIZE * sizeof(int), cudaMemcpyDeviceToHost);
```

```
printf("{1,2,3,4,5} + {10,20,30,40,50} = {%d,%d,%d,%d,%d}\n",  
       c[0], c[1], c[2], c[3], c[4]);
```

```
cudaFree(dev_c);  
cudaFree(dev_a);  
cudaFree(dev_b);  
return 0;
```

```
}
```



Example code (Assignment)

■ Result?

– Quick assignment

- Write the vector addition code with the SIZE of 5
- The assignment code should run $d[i] = a[i] + b[i] + c[i]$.
(NOT the same as the example code)
- The values of source arrays (a,b, and c) can be initialized with random values.
- You should capture your source code and the results.



Summary

- Kernel concept
- CUDA programming model
- CUDA function declarations
- Vector addition example
- CPU implementation
 - For loop
- Kernel concept
 - Function, as loop-body
- CPU kernel implementation
- CUDA implementation
 - Multiple thread launch
- Kernel launch
 - Example code and **Assignment!**



Next step?

- CUDA kernel implementation
 - 2D matrix examples?



Thank you

Any questions?

E-mail: yhgong@kw.ac.kr

Lab: <https://sites.google.com/view/yhgong/>



광운대학교
KwangWoon University

